

LuMamba: one backbone, two-pre-training objectives

Topology-invariant across 16-26 channels • $377\times$ fewer FLOPS than LaBraM • $12\times$ longer sequences

